

Yiwei Chen

 Github |  Homepage |  LinkedIn |  Google Scholar

EDUCATION

Michigan State University

East Lansing, USA

- **Ph.D. in Computer Science**

Aug. 2024 - Present

Advisor: Prof. [Sijia Liu](#)

Xi'an Jiaotong University

Xi'an, China

- **M.S. in Computer Science and Technology**

Sept. 2021 - Jun. 2024

- **B.Eng. in Automation**

Sept. 2017 - Jun. 2021

Admitted via the [Special Class for the Gifted Young](#), a national early-entrance honors program

RESEARCH INTERESTS

- Reasoning, Post-training, LLMs, MLLMs, Agentic System, GUI Agent
- AI Safety, Alignment, Trustworthy Algorithms, Interpretability, Machine Unlearning

PUBLICATIONS & MANUSCRIPTS

(* means Equal Contribution)

- [1] [Y. Chen](#), Y. Yao, Y. Zhang, B. Shen, G. Liu, S. Liu. "Safety Mirage: How Spurious Correlations Undermine VLM Safety Fine-Tuning and Can Be Mitigated by Machine Unlearning." *ICLR*, 2026. [\[Paper\]](#)
- [2] [Y. Chen](#), S. Pal, Y. Zhang, Q. Qu, S. Liu. "Unlearning Isn't Invisible: Detecting Unlearning Traces in LLMs from Model Outputs." *ICLR*, 2026. [\[Paper\]](#)
- [3] [Y. Chen](#), B. Shang, S. Liu. "Who Built This Model? Tracing LLM Lineage via Spectral Fingerprints in Weight Space." *COLM*, 2026. [\[Paper\]](#)
- [4] [Y. Chen](#), L. Li, K. Cheung, V. Parla, G. Sundaram. "Data-Centric Benchmarking of Exploit Generation in LLMs: Understanding the Impact of Fine-Tuning." Manuscript, 2026. [\[Paper\]](#)
- [5] [Y. Chen](#)*, B. Shang*, Y. Zhang, B. Shen, S. Liu. "Forgetting to Forget: Attention Sink as A Gateway for Backdooring LLM Unlearning." *COLM*, 2026. [\[Paper\]](#)
- [6] [Y. Chen](#)*, Z. Zhang*, W. Zhang, C. Yan, Q. Zheng, Q. Wang, W. Chen. "Tile Classification Based Viewport Prediction with Multi-modal Fusion Transformer." *ACM-MM*, 2023. [\[Paper\]](#)
- [7] Y. Zhang, C. Wang, [Y. Chen](#), C. Fan, J. Jia, S. Liu. The Fragile Truth of Saliency: Improving LLM Input Attribution via Attention Bias Optimization. *NeurIPS Spotlight*, 2025. [\[Paper\]](#)
- [8] H. Reisizadeh, J. Ruan, [Y. Chen](#), S. Pal, S. Liu, M. Hong. Leak@k: Unlearning Does Not Make LLMs Forget Under Probabilistic Decoding. *ICML*, 2026. [\[Paper\]](#)

INTERNSHIPS

Amazon, Applied Scientist Intern

May 2026 - Aug. 2026

Research Project: Memory-Augmented World Models for Mobile GUI Agents

- Identified the loop-closure problem in GUI world models, where one-step predictors regenerate a revisited page instead of restoring its earlier state, and extended the task to trajectory-level prediction with loop-closure evaluation.
- Built an automated annotation pipeline that adds page identity and revisit structure to three GUI benchmarks, and a training pipeline for a latent transition memory around a frozen VLM, outper-

forming existing GUI world models on consistency. Manuscript in preparation for CVPR 2027.

Cisco, Research Intern

May 2025 - Dec. 2025

Research Project: LLMs for CVE Exploit Generation in Cybersecurity [4]

- Built a data-centric benchmark and fine-tuned Qwen3-8B for CVE-conditioned exploit generation, achieving a +42.5% improvement over the base model, comparable with closed-source LLMs.

RESEARCH EXPERIENCE

OPTML, Michigan State University

East Lansing, US

Research Assistant. Advisor: Prof. [Sijia Liu](#)

Aug. 2024 -

- *Multi-Agent Systems*: Formulated multi-agent LLM training as consensus optimization over discrete nodes, and proposed an ADMM framework that synchronizes only low-rank adapter states while provably matching centralized training. Manuscript in preparation for ICLR 2027.
- *GUI Agents*: Studied GUI-agent systems with a focus on world modeling, state prediction, action planning, and memory-based agent decision making in interactive environments.
- *Trustworthy Foundation Models* [1], [2], [3], [5]: Studied how to make foundation models more reliable, accountable, and robust after training, focusing on safety alignment, knowledge removal, model provenance, and hidden vulnerabilities that may persist or emerge during fine-tuning and unlearning.

AWARDS & HONORS

SCHOLARSHIPS

- ICLR 2026 Travel Award Apr. 2026
- Michigan State University Travel Award Oct. 2025
- Outstanding Master Graduate, Xi'an Jiaotong University (Top 5% of all graduates) Jun. 2024
- Outstanding Bachelor Graduate, Xi'an Jiaotong University (Top 5% of all graduates) Jun. 2021
- First Prize Scholarship, Xi'an Jiaotong University (Top 10%) Nov. 2016-2019, 2022-2023

COMPETITIONS

- Meritorious Winner in Mathematical Contest in Modelling (top 8% in 25370 teams) Mar. 2019
- The First Prize in Shaanxi of China Undergraduate Mathematical Contest in Modeling Oct. 2018

ACADEMIC SERVICE

- **Conference Reviewer:** ICLR, NeurIPS, ICASSP, ACM-MM, IJCNN
- **Journal Reviewer:** Journal of Machine Learning Research (JMLR)

LANGUAGE & SKILLS

Language	Mandarin (Native), English (Professional)
Programming	Python, C/C++, Bash, LaTeX, MATLAB, HTML/CSS
Model Frameworks	LLMs (e.g., Llama, Qwen), MLLMs (e.g., LLaVA, Qwen-VL)
Libraries / Softwares	Pytorch, Tensorflow, Deepspeed, Huggingface, OpenCV
Developer Tools	Git, Docker, Vim, VSCode, Cursor, Claude Code, PyCharm